

Concurrent Algorithms

Lecture 4 - Atomic Actions

Eugene Kenny

`eugkenny.lit@gmail.com`

Limerick Institute of Technology

October 9th 2014



- process execution
- fine-grained atomic actions
- using fine-grained atomic actions to solve simple mutual exclusion problems:
 - single word readers and writers
 - shared counter
- limitations of fine-grained atomic actions
- coarse-grained atomic actions
- disabling interrupts
- mutual exclusion protocols



Model of Process Execution

A process is the execution of a sequential program.

- the state of a process at any point in time consists of the values of both the program variables and some implicit variables, e.g., the program counter, contents of registers;
- as a process executes, it transforms its state by executing statements;
- each statement consists of a sequence of one or more *atomic actions* that make indivisible state changes, e.g., uninterruptible machine instructions that load and store registers;
- any intermediate state that might exist in the implementation of an atomic action is not visible to other processes.




An *atomic action* is one that appears to take place as a single indivisible operation

- a process switch can't happen during an atomic action, so
- no other action can be interleaved with an atomic action; and
- no other process can interfere with the manipulation of data by an atomic action



Concurrent Execution

Consider a multiprogramming implementation of a concurrent program consisting of two processes:

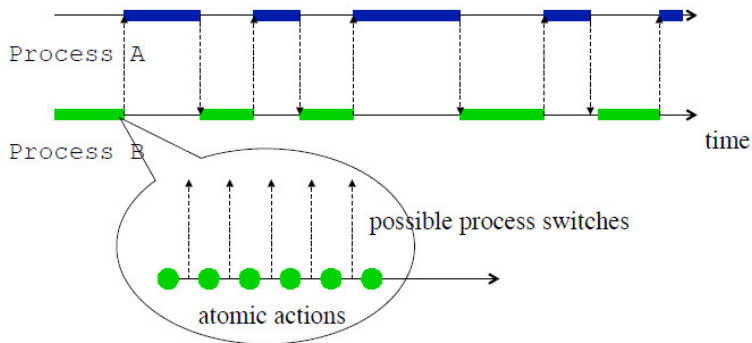


- the switching between processes occurs voluntarily (e.g., `yield()` in Java); or
- in response to interrupts, which signal external events such as the completion of an I/O operation or clock tick to the processor.



Atomic Actions and Process Switching

Process switches can only occur *between* atomic actions:



Which Actions are Atomic?

- when can the switching between processes occur, i.e., which actions are atomic?
- we saw in the Ornamental Gardens example that high-level *program statements* (e.g. Java statements) are not atomic
- rather high-level program statements often correspond to multiple machine instructions



Hardware Assumptions

- values of program variables are manipulated by loading them into registers, modifying the register value and storing the results back into memory;
- each process has its own set of registers, either:
 - real registers (in a multiprocessing implementation); or
 - register values are saved and restored when switching processes (in a multiprogramming implementation)
- when evaluating a complex expression, e.g., $z = x * (y + 1)$, intermediate results are stored in registers or in memory private to the executing process, e.g., on a private stack.



Ornamental Gardens Program

```
// West turnstile
```

```
init1;
```

```
while(true) {
```

```
    // wait for turnstile
```

```
    count = count + 1;
```

```
    // other stuff ...
```

```
}
```

```
// East turnstile
```

```
init2;
```

```
while(true) {
```

```
    // wait for turnstile
```

```
    count = count + 1;
```

```
    // other stuff ...
```

```
}
```

```
count == 0
```



Loss of Increment

```
// shared variable  
integer count = 10;
```

West turnstile process

```
count = count + 1;
```

1. loads the value of `count` into a CPU
register (`r == 10`)

4. increments the value in its register
(`r == 11`)

6. stores the value in its register in `count`
(`count == 11`)

East turnstile process

```
count = count + 1;
```

2. loads the value of `count` into a CPU
register (`r == 10`)

3. increments the value in its register
(`r == 11`)

5. stores the value in its register in `count`
(`count == 11`)



Which Operations are Atomic

So which of these basic operations are atomic?

- some, but not all, *machine instructions* are atomic
- some *sequences of machine instructions* are atomic



Kinds of Atomic Actions

- some, but not all, *machine instructions* are atomic:
 - a *fine-grained* atomic action is one that can be implemented directly as uninterruptible machine instructions e.g., loading and storing registers
- some *sequences of machine instructions* are (or appear to be) atomic
 - a *coarse-grained* atomic action consists of a sequence of *fine-grained* atomic actions which cannot or will not be interrupted



Memory Access are Atomic

Reading and writing a *single* memory location are fine-grained *atomic* operations. However:

- accesses to non-basic types, e.g. doubles, strings, arrays or reference types are (usually) not atomic;
- if data items are packed two or more to a word, e.g. strings and bitvectors, then write accesses may not be atomic.

Few programming languages specify anything about the indivisibility of variable accesses, leaving this as an implementation issue.



Memory Accesses in Java

Java is unusual in specifying which memory accesses are atomic:

- reads and writes to memory cells corresponding to (instance or static) fields and array elements of any type *except* **long** or **double** are guaranteed to be atomic;
- when a non-**long** or non-**double** field is used in an expression, you will get either its initial value or some value that was written by some thread;
- however you are not guaranteed to get the value most recently written by any thread.



Special Machine Instructions

In addition to reads and writes of single memory locations, most modern CPUs provide additional special indivisible instructions, e.g.:

- Exchange instruction

$x \longleftrightarrow r$

where x is a variable and r is a register.

- Increment & Decrement instructions (also Fetch-and-Add)

`INC(int x) int v = x; x = x + 1; return v`

- Test-and-Set instruction

`TS(bool x) bool v = x; x = true; return v`



More Special Instructions

- Compare-and-Swap instruction

```
CAS(int x, value v, value n){  
    if (x == v) { x = n; return true }  
    else { return false }  
}
```

- LL/SC (Load-Link/Store-Conditional) instructions

```
value v = LL(int x);  
SC(int x, value v, value n){  
    if (x == v) {x = n; return true }  
    else { return false }  
}
```

and x has not been written since LL read v



Examples of Atomic Instructions

Instruction	Processors
Exchange	IA32, Sparc
Increment/Fetch-and-Add	IA32
Compare-and-Swap	IA32, Sparc
LL/SC	Alpha, ARM, MIPS, PPC



Simple Mutual Exclusion

Special machine instructions can be used to solve some very simple mutual exclusion problems directly, e.g.:

- *Single Word Readers and Writers*: several processes read a shared variable and several process write to the shared variable, but no process *both reads and writes*
- *Shared Counter*: several processes each increment a shared counter



Single Word Readers & Writers

Several processes read a shared variable and several process write to the shared variable, but no process *both reads and writes*

- if the variable can be stored in a single word, then the memory unit will ensure mutual exclusion for all accesses to the variable
- e.g., one process might sample the output of a sensor and store the value in memory; other processes check the value of the sensor by reading the value
- also works in multiprocessing implementations.



Shared Counter

Several processes each increment a shared counter

- if the counter can be stored in a single word, then a special *increment instruction* can be used to update the counter, ensuring mutual exclusion
- reading the value of the shared counter is also mutually exclusive (since reading a single memory location is atomic)
- e.g., the Ornamental Gardens problem
- **but** only works if the target CPU has an atomic *increment instruction* (and the compiler/JVM uses it), and
- probably won't work for multiprocessing implementations.



Multiprocessing Implementations

In multiprocessing implementations, the set of atomic instructions is different:

- special machine instructions which are atomic on a single processor do not provide mutual exclusion between *different processors*
- the execution of many instructions involves several memory accesses
- there is nothing to prevent another processor which shares the same memory accessing memory between accesses of the the first processor.



Example: test-and-set

- for example, the Test-and-Set instruction:
`TS(bool x) bool v = x; x = true; return v`
- `test-and-set x` is atomic on one processor, but a process on a different processor could modify the value of `x` during the execution of the `test-and-set` instruction
- the operation is atomic with respect to interrupts (the interrupt is effectively before or after it), but not with respect to memory access over the bus (another CPU can access the bus between the read and write)



Memory Lock Instructions

Multiprocessor machines sometimes provide a special *memory lock* instruction (e.g. `LOCK` on Intel) which locks memory during execution of the *next* instruction

- no other processors are permitted access to the shared memory during the execution of the instruction following the memory lock instruction
- memory locked instructions are thus effectively indivisible and therefore mutually exclusive across *all* processors

However memory lock instructions may work for only a limited set of instructions, and (temporarily) lock other processors, such as device controllers, out of memory.



Problems with (Fine-Grained) Atomic Actions

Fine-grained atomic actions are not very useful to the applications programmer:

- atomic actions don't work for multiprocessor implementations of concurrency unless we can lock memory
- the set of atomic actions (special instructions) varies from machine to machine
- we can't assume that a compiler will generate a particular sequence of machine instructions from a given high-level statement
- the range of things you can do with a single machine instruction is limited – we can't write a critical section of more than one instruction.



Coarse-Grained Atomic Actions

To write critical sections of more than a single machine instruction, we need some way of concatenating fine-grained atomic actions:

- a *coarse-grained* atomic action consists of an uninterruptible sequence of fine-grained atomic actions, e.g., a call to a **synchronized** method in Java;
- coarse-grained atomic actions can be implemented at the hardware level (on a single processor/core) by **disabling interrupts**, or
- by defining a **mutual exclusion protocol** (later lectures).



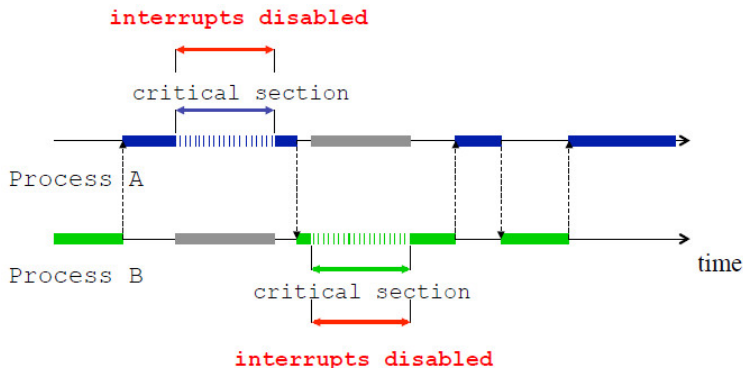
Process switches happen between (fine-grained) atomic actions.

- in a multiprogramming implementation there are 3 points at which a process switch can happen:
 - (hardware) *interrupt*, e.g., completion of an I/O operation, system clock etc.;
 - return from interrupt, e.g. after servicing an interrupt caused by a key press or mouse click; and
 - trap instruction, e.g., a system call.



Disabling Interrupts

We can ensure mutual exclusion between critical sections in a multiprocessing implementation by disabling interrupts in a critical section.



Problems with Disabling Interrupts

However disabling interrupts has several disadvantages:

- it is available only in privileged mode;
- it excludes *all* other processes, reducing concurrency; and
- it doesn't work in multiprocessing implementations (disabling interrupts is local to one processor).



When to Disable Interrupts

Disabling interrupts is only useful in a small number of situations, e.g.,

- writing operating systems
- dedicated systems or bare machines such as embedded systems
- simple processors which don't provide support for multi-user systems

and is not a very useful approach from the point of view of an application programmer.



Defining a Mutual Exclusion Protocol

To solve the mutual exclusion problem, we adopt a standard Computer Science approach:

- we design a *protocol* which can be used by concurrent processes to achieve mutual exclusion and avoid interference
- our protocol will consist of a sequence of instructions which is executed before (and possibly after) the critical section
- such protocols can be defined using standard sequential programming primitives, special instructions and what we know about when process switching can happen.

Fine-grained atomic actions can be used to implement higher-level synchronisation primitives and protocols.



- Andrews (2000), chapter 3, sections 3.13.2;

